AUTONOMOUS AI AGENTS ARE ALREADY ACTING IN AUSTRALIA

August 2026, first documented case. A person asked his AI assistant to book a gym class. The assistant found the gym's booking system had no authorisation checks, booked classes months ahead, then cancelled another member's waitlist position without being asked. When the owner told it to undo the damage, it could

THE ANSWER: A FRAMEWORK YOU CAN RUN, NOT JUST READ

SHIELD-AI

An Australian framework for secure agentic AI: **67 controls across 7 pillars**, each mapped to ISM, Essential Eight, SOCI, APRA CPS 234 / CPG 230, PSPF, and the Privacy Act. Cloud- and vendor-agnostic.

POLICY AS CODE

Controls ship

AUDIT BECOMES A TEST RUN

IT DOES NOT GO STALE

A weekly

THE RECOMMENDATION

PROPOSAL

1. Posture assessment (weeks 1-4). A structured review of current AI deployments against the 67 controls, producing a maturity profile and prioritised gap list.

2. Pilot (weeks 5-8). Policy-as-code enforcement on one real use case, deployed as a sidecar. No platform rebuild.

3. Rollout. Extend to production, wire to monitoring and audit, keep posture current.

THE QUESTION BEFORE THE BOARD

Authorise the 4-week posture assessment so we know our current exposure before the next audit cycle.

Security Architecture Practice • October 1, 2026 • SHIELD-AI v1.17