SHIELD-AI™

Security and Holistic Integrity for Enterprise-Level Deployment of AI

Technical Deep Dive -- Framework Architecture, Controls & Implementation

Security Architecture Practice * v1.7 * August 2026

AGENDA

1. Architecture Overview

- Reference architecture -- control stack
- Threat landscape & attack taxonomy

2. Seven Pillars Deep Dive

- MS * IAM * DS * TP * MIR * GRM * EE
- Controls, decision trees, mappings

3. Regulatory & Standards

- AU frameworks, OWASP, NIST, EU AI Act, ISO 42001

4. Emerging Trends & Protocol Security

- Multi-agent, MCP, persistent memory, code-gen agents

5. Delivery & Operations

- Implementation methodology, continuous compliance
- Deployment topologies, reference implementation, deployable components

6. Reference & Appendices

REFERENCE ARCHITECTURE -- SHIELD-AI™ CONTROL STACK

This is the anchor for every pillar. Each numbered area maps to a SHIELD-AI™ pillar.



AGENTIC AI THREAT LANDSCAPE

Seven layers of attack surface -- controls must cover all of them

🧠 MODEL LAYER

Prompt injection, jailbreaking, model inversion, data extraction from responses, backdoor triggers

📊 DATA LAYER

Cross-session data aggregation, RAG pipeline poisoning, PII leakage in tool calls

🔧 TOOL LAYER

Tool output injection, malicious tool registration, API abuse, tool chain poisoning

🕹️ PIPELINE LAYER

CI/CD compromise, model update substitution, dependency chain poisoning

🔒 IDENTITY LAYER

🏛 GOVERNANCE LAYER

WHY EXISTING FRAMEWORKS DON'T COVER AGENTIC AI

ASD ISM (1,081 CONTROLS)

Zero AI-specific controls. March 2026 "Guidelines for secure AI system development" is introductory, not a control framework.

ESSENTIAL EIGHT

Necessary but insufficient -- doesn't address prompt injection,

OWASP LLM TOP 10

Threat taxonomy, not a framework with maturity levels, implementation guidance, or regulatory mapping.

EU AI ACT

Risk classification + obligations. No technical controls defined.

ISO/IEC 42001:2025

SHIELD-AI FRAMEWORK -- 7 PILLARS X 51 CONTROLS

P1

MODEL SECURITY

Weight protection, training integrity, supply chain, adversarial robustness, behaviour validation, inventory, context obfuscation, MCP controls

P2

IDENTITY & ACCESS

Agent identity, scope-bound delegation, human oversight, credential isolation, JIT access, cross-session, delegation chains, credential scoping, multi-tenant isolation

P3

DATA SECURITY

Classification, inference protection, memory sanitisation, minimisation, training governance, output validation

P4

TOOL & PIPELINE

Tool registry, access control, API gateway, output validation, pipeline integrity, dependency scanning, chain analysis,

P5

MONITORING & IR

Activity logging, anomaly detection, prompt audit, IR playbooks, reporting integration, forensic readiness

P6

GOVERNANCE & RISK

Governance body, risk register, ethical use, third-party risk, regulatory mapping, board reporting

SHIELD-AI VS "WE ALREADY USE NIST + OWASP + ISO 42001"

The most common technical objection, answered directly

Dimension	NIST AI RMF + OWASP + ISO 42001	SHIELD-AI
Framework count	3 separate, each with different terminology	1 unified framework
Maturity model	None -- no "where are we → where next"	4 levels (L1-L4) with per-control progression
AU regulatory mapping	None -- must build custom mapping	Full ISM/E8/SOCI/CPS 234/PSPF/Privacy Act
Agentic-AI specificity	Partial -- NIST AI RMF covers general AI risk	Full -- 7 pillars designed for autonomous agents
	No -- process/procedure	

PILLAR 1: MODEL SECURITY (MS)

MS-01 **Weight protection** -- encrypt at rest/transit, HSM-backed keys, remote attestation for edge — ISM

MS-02 **Training data integrity** -- provenance tracking, sanitisation, poison detection pipeline — ISM, LLM04

MS-03 **Supply chain** -- model provenance, ML SBOM, signed checkpoints — ISM, LLM10

MS-04 **Adversarial robustness** — LLM01, NIST

DECISION TREE: CHOOSING YOUR APPROACH

💡 **How to approach model security:**

Q: Do you use third-party or self-hosted models?

> Third-party → focus on MS-04 (prompt injection testing), MS-05 (output guardrails)
>
> Self-hosted → add MS-01 (weight encryption), MS-02 (training integrity)

Q: How many models in production?

> <5 → Manual inventory (MS-06

PILLAR 2: IDENTITY & ACCESS (IAM)

IAM-01 **Agent identity** -- unique service identities per agent, not human credential sharing ISM E8

IAM-02 **Scope-bound delegation** -- tools scoped to agent function, no wildcards, no self-escalation LLM08

IAM-03 **Human oversight** -- approval gates for high-risk actions, break-glass with audit trail

IAM-04 **Credential isolation** -- agent

DECISION TREE: CHOOSING YOUR APPROACH

💡 **Choosing your agent identity approach:**

Q: How many agents?

> <5 → OPA sidecar per agent (simplest path)
>
> 5-50 → Central OPA gateway (single policy enforcement point)
>
> 50+ → OPA + SPIFFE/SPIRE for workload identity (scales)

Q: Multi-agent architecture?

PILLAR 3: DATA SECURITY (DS)

DS-01 **Data classification** -- automated tagging per PSPF/ISM at ingestion — ISM

DS-02 **Inference protection** -- TLS 1.3, end-to-end encryption of prompts/responses

DS-03 **Memory sanitisation** -- no sensitive data across sessions; retention policies; automated purging — LLM06

DS-04 **Data minimisation** -- field-level filtering in retrieval layer, no bulk data

DECISION TREE: CHOOSING YOUR APPROACH

💡 **How to approach data security:**

Q: Does your agent use persistent memory?

> Yes → DS-03 is highest priority: retention policies, automated purging, per-agent scoping
>
> No → Defer DS-03, focus on DS-01 (classification) and DS-04 (minimisation)

Q: RAG pipeline?

> Yes → Implement ABAC at retrieval

PILLAR 4: TOOL & PIPELINE SECURITY (TP)

TP-01 **Tool registry** -- all tools registered, versioned, security-approved

TP-02 **Tool access control** -- least-privilege per agent, tool authenticates agent

TP-03 **API gateway** -- rate limiting, schema validation, audit logging on agent-facing APIs

TP-04 **Tool output validation** -- structural validation, sanitisation before agent context LLM02

DECISION TREE: CHOOSING YOUR APPROACH

💡 **Choosing your deployment topology:**

> <10 agents → Embedded library (simplest, framework-native)
>
> 10-100 agents → Sidecar proxy (balanced isolation vs overhead)
>
> 100+ / multi-agent → Central gateway for external, sidecars for cross-agent

PILLAR 5: MONITORING & IR (MIR)

MIR-01 **Activity logging** -- every tool call, data access, decision; agent ID + session + correlation ID

MIR-02 **Anomaly detection** -- per-agent and per-class baselines; deviation alerts

MIR-03 **Prompt audit log** -- sanitised prompt/response pairs, recoverable for IR

MIR-04 **IR playbooks** -- prompt injection containment, model rollback, tool revocation

MIR-05 **Mandatory reporting** SOCI CPS 234

DECISION TREE: CHOOSING YOUR APPROACH

💡 **How to build agent telemetry:**

Q: Existing SIEM?

> Splunk/SOAR/ELK → Deploy agent telemetry as structured JSON events via OTLP
>
> None → Start with cloud logging (CloudWatch, Azure Monitor, GCP Cloud Logging)

Q: Multi-agent architecture?

> Yes → Distributed tracing header across all agents; single audit trail

PILLAR 6: GOVERNANCE & RISK (GRM)

GRM-01 **AI governance body** -- CISO, CRO, GC, data owner, AI lead with charter CPS 234

GRM-02 **AI risk register** -- distinct taxonomy for model/tool/data/identity/pipeline risks

GRM-03 **Ethical use policy** -- acceptable use, human-in-the-loop by risk tier, transparency

GRM-04 **Third-party AI risk** -- vendor assessment, contractual controls, right-to-audit LLM10

DECISION TREE: CHOOSING YOUR APPROACH

💡 **How to build AI governance:**

Q: Existing risk/compliance function?

> Mature (APRA-regulated) → GRM-01 (chartered committee), GRM-02 (risk register), GRM-05 (automated mapping)
>
> Developing → Start with GRM-06 (dashboard -- shows posture without formal governance)

Q: Using external model providers?

PILLAR 7: EXECUTION ENVIRONMENT (EE)

E8

EE-01 **Boundary enforcement** -- container/VM isolation, no shared mounts/namespaces, immutable infra, fail-closed

EE-02 **Infrastructure API control** -- no implicit access to EC2/IAM/Lambda/S3 control planes

EE-03 **Resource limiting** -- CPU/memory/network/storage/API rate bounds; exceed = alert or terminate

EE-04 **Evaluation sandbox integrity** -- eval isolated from prod, no network path, no prod

DECISION TREE: CHOOSING YOUR APPROACH

💡 **How to build the execution boundary:**

Q: Container or VM?

> <10 agents → Per-agent containers with hardened runtime profiles
>
> Untrusted/unverified code → MicroVM or VM-level isolation (no virtiofs/9p passthrough)

Q: Does the agent execute LLM-generated code?

> Yes → EE-06 is mandatory; runtime

MATURITY MODEL -- LEVEL-UP GUIDANCE

Area	L1 Initial	L2 Managed	L3 Defined	L4 Optimised	How to Le Up (exam IAM)
MS	Basic inventory	Supply chain verification	Adversarial testing programme	Automated red-teaming	L1→L2: Sig ML SBOM L2→L3: Deploy CI/ prompt injection tests. L3→ Continuou red-teami
		Per-agent			L1→L2: Ea agent gets identity. L2 >L3: OPA

MCP PROTOCOL & AGENT-TOOL INTERFACE SECURITY

✅ SECURITY BENEFITS OF MCP

Standardised interfaces enable standardised controls. MCP defines tool discovery schemas, typed tool interfaces (JSON Schema), and structured error responses -- making generic tool security proxies possible.

SHIELD-AI MCP CONTROLS

TP-03 extended: MCP gateway enforces auth, schema validation, rate limiting at protocol layer. **TP-04 extended:** MCP response validation against known-good schemas. **MS-04 extended:** MCP protocol-layer adversarial testing.

⚠️ MCP ATTACK SURFACE

MCP GATEWAY ARCHITECTURE

EMERGING: MULTI-AGENT SYSTEMS & AGENT SWARMS

ARCHITECTURE CHANGES

- **Delegate** tasks to other agents -- delegation chains
- **Share context** -- one agent's output is another's input
- **Coordinate** -- interdependent multi-agent workflows

BLAST RADIUS MULTIPLIER

A single compromised agent poisons every downstream agent's context. Prompt injection

IAM-02: BOUNDED DELEGATION CHAIN

Agent-to-agent must not expand original scope. Each hop respects bounds -- no escalation across delegation chains.

MIR-01: DISTRIBUTED TRACING

Single user request → chain of agent calls reconstructable as one

EMERGING: PERSISTENT AGENT MEMORY

DATA AGGREGATION RISK

- Agent remembers User A's financial details
- User B queries in Session 2 -- cross-reference possible
- Memory embedding spaces encode sensitive correlations

ATTACK VECTORS

Memory-read: compromised session reads another session's data. **Memory poisoning:** crafted inputs trigger unintended

DS-03: MEMORY SANITISATION (HIGHEST PRIORITY)

Retention policies with automated purging. Cross-session aggregation detection. Per-agent memory scoping. Encryption at rest.

DS-04: MEMORY MINIMISATION

Only essential data remembered. Memory scoped per-agent and per-task -- not shared across

EMERGING: AUTONOMOUS CODE GENERATION

The highest-risk agent capability -- treat accordingly

Level	Capability	Human Approval
1	Generate SQL queries	Always
2	Generate shell commands	Always
3	Generate config files	Always
4	Self-heal -- patch own code	Always + sandbox
5	Write and deploy microservices	Multi-party approval

TP-01: CODE EXECUTION AS TOOL CLASS

Distinct class with highest security. Must be explicitly enabled per agent -- never default-on.

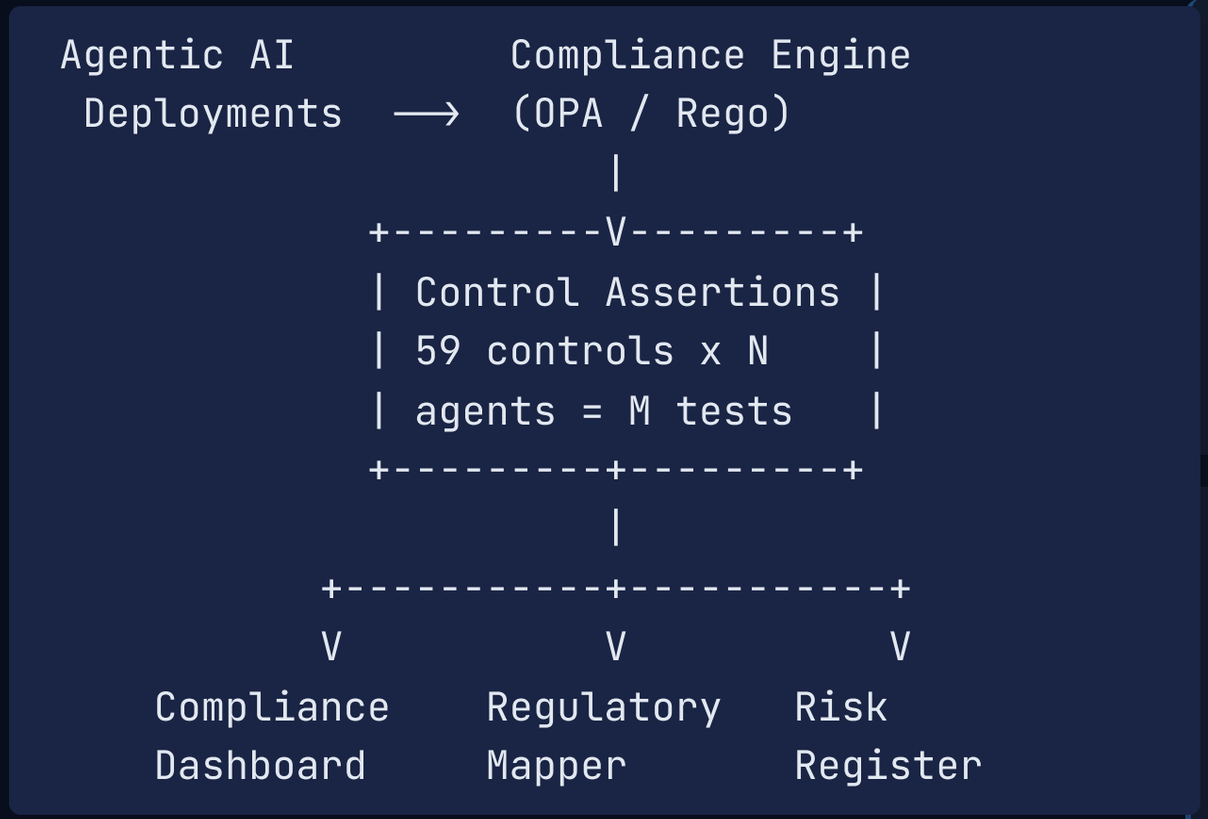
TP-04: CODE OUTPUT VALIDATION

Sandbox-executed and output-

AU REGULATORY MAPPING

SHIELD-AI	ISM	E8	SOCI Act	CPS 234	Privacy Act
MS-01	Cryptography	-	CIRMP s.4(a)	s.3(b)	APP 11
IAM-01	Sys Mgmt	Restrict Priv	-	s.3(b)	-
IAM-02	Sys Mgmt	-	CIRMP s.4(b)	s.3(a)	-
DS-01	Data Transfers	-	-	s.3(a)	APP 5,6
DS-03	Media	-	-	-	APP 11
TP-01	Gateways	Patch Apps	CIRMP s.4(c)	s.3(b)	-
MIR-01	Monitoring	-	s.30B(5)(a)	s.3(c)	NDB
MIR-05	Incidents	-	s.30B(5)(c)	s.3(c)	NDB
GRM-01	Roles	-	CIRMP s.4(d)	s.2	APP 1

CONTINUOUS COMPLIANCE AUTOMATION



1. CONTROL AS CODE

Every SHIELD-AI control = machine-readable assertion (Rego policy, CI/CD test, etc.)

2. CONTINUOUS TESTING

Critical controls: every 15 min. Config controls: daily. Dashboard with trend data.

3. REGULATORY CHANGE DETECTION

DEPLOYMENT TOPOLOGIES -- DECISION TREE

<10 AGENTS

Embedded library -- security controls via agent framework SDK/plugin. Low latency, simple config. *Best for: pilots, single-agent PoCs.*

10-100 AGENTS

Sidecar proxy -- per-agent sidecar (OPA + tool gateway + telemetry). *Best for: K8s deployments, existing service mesh.*

100+ / MULTI-AGENT

Central gateway + sidecar audit -- gateway for external tools, sidecars for cross-agent monitoring. *Best for: enterprise, regulated, multi-agent systems.*

PROS

No infra dependency, consistent with framework API

PROS

Isolation per agent, no SPOF, consistent with sidecar mesh

PROS

Single policy point, central

BUILD-VS-BUY ANALYSIS

Factor	Build In-House	Adopt SHIELD-AI
Time to capability	12-18 months	4-16 weeks
Hiring	3-5 specialists (6-9mo to hire)	None -- implementation team delivers
Framework depth	Inexperienced, likely gaps	59 controls, 12+ framework mappings
Regulatory confidence	Low -- untested	High -- auditable mappings
Incident learnings	Reactive -- each incident a new lesson	Built-in -- 36 documented incidents drive the control set

WHY ADOPT

Every control is a documented, testable assertion with implementation guidance -- not a policy binder. Deployable components are specified in the implementation architecture document.

WHY BUILD FAILS

The threat landscape moves monthly (CrewAI, Codex CLI, CVE-

MATURITY LEVELS -- TECHNICAL SUMMARY

Capability	L1 Initial	L2 Managed	L3 Defined	L4 Optimised
Assessment scope	Single pilot / PoC	Priority AI systems	Full agentic AI estate	Multi-entity / cross-border
Control implementation	Basic model controls	Documented controls, per-agent accounts	Automated L3 deployment	Custom control development
Technical tooling	Manual inventory	Model registry, basic gateway	OPA policies, tool gateway, SIEM dashboards	+ Threat intel, red-teaming platform

KEY TECHNICAL TERMS

Term	Definition	Control
MCP	Standardised protocol for agent-tool discovery and invocation	TP-03, TP-04
A2A Protocol	Standard for direct inter-agent communication and delegation	IAM-02, MIR-01
Policy-as-Code	Security policies as machine-readable, testable assertions (OPA/Rego)	GRM-05, IAM-02
Tool Output Injection	Compromised tool returns crafted data to manipulate agent behaviour	TP-04, DS-06
Agent SBOM	Software bill of materials: model, tools, libraries, pipeline, runtime	MS-03, TP-06
RAG Security Layer	ABAC at retrieval + field-level redaction before model context	DS-01, DS-04
Agent Telemetry	Structured event stream: tool calls, data access, decisions, errors	MIR-01, MIR-02
Bounded Delegation	Each agent-to-agent hop respects original scope -- no escalation	IAM-02
Continuous Compliance	Control verification runs continuously, not point-in-time	GRM-05, GRM-06
MCP Gateway	Security proxy for MCP protocol: auth, schema, audit, rate-limiting	TP-03

TRY IT TODAY -- IMPLEMENTATION ARCHITECTURE

💻 IMPLEMENTATION ARCHITECTURE

framework/shield-ai-implementation-architecture.md -- deployable components, reference architecture, full SVG diagram set, component-to-control mapping.

📋 POLICY-AS-CODE

policies/ -- 36 Rego stubs covering MS, IAM, DS, TP, MIR, GRM. Wire them into the OPA decision points from the architecture document.

🚀 ASSESSMENT TOOLKIT

The assessment guide scores all 59 controls per pillar. Control schema (JSON) is machine-readable for tooling.

QUESTIONS & DISCUSSION

SHIELD-AI v1.7 -- August 2026

Security Architecture Practice

Full framework document • Implementation architecture (deployable components + diagrams) • Gap analysis • Assessment guide