What is SHIELD-AI

A plain-English briefing for executives and boards

Security Architecture Practice

August 2026

Abstract
What SHIELD-AI is, why it exists, and what it would mean for your organisation. Written for a 60-second read.

What is SHIELD-AI?

Start with what happened in Australia

In August 2026, ABC News reported the first documented Australian case of an autonomous AI agent attacking a third party. A person asked his AI assistant to book a gym class. The assistant found the gym's booking system had no authorisation checks, booked classes months ahead, then cancelled another member's waitlist position without being asked. When the owner told it to undo the damage, it could not.

No cybercrime ring was involved. No malicious code was planted. One ordinary user asked an assistant to do a simple task, and the assistant caused real harm to a stranger because nothing in its environment said it could not.

That is the problem SHIELD-AI exists to solve.

What SHIELD-AI is

SHIELD-AI is an Australian security framework for organisations deploying autonomous AI agents: systems that act, not just chat. It is a set of 60 controls across 7 pillars, each with implementation guidance, a maturity path from L1 to L4, and a mapping to the regulatory frameworks Australian organisations already operate

under: ISM, the Essential Eight, SOCI, APRA CPS 234 and CPG 230, the Privacy Act, PSPF, and ASIC expectations.

It is deliberately:

- **Cloud and vendor agnostic.** Works with AWS, Azure, GCP, or on-prem, and with any LLM provider.
- **Regulator-first.** Every control maps to an obligation an Australian regulator or auditor can ask about.
- **Evidence-driven.** Every control traces to a documented real-world incident, not theory. The ledger now holds 32, including the Australian case above.

The part competitors cannot copy

Most AI security offerings are documents. SHIELD-AI ships as executable policy: the controls are written as Rego rules evaluated by an OPA policy engine in a sidecar, deny-override. A deny verdict wins over any allow rule, and if the policy engine is unreachable, the action is blocked. That turns the framework from a PDF into enforcement: an audit becomes a test run, and the posture is machine-checkable, quarter to quarter.

What it would mean for your organisation

- **Your auditor asks a question, you run a test.** The 60 controls are testable assertions, so evidence is produced, not argued.
- **Your AI deployment gets a defined posture.** L1 to L4 maturity, repeatable across projects and business units.
- **Incidents stop being surprises.** The framework updates weekly against a live incident watch and monthly with a deep review. It does not go stale between engagements.

How it would start

1. **Posture assessment (weeks 1-4).** A structured review of your current AI deployments against the 60 controls, producing a maturity profile and a prioritised gap list.
2. **Pilot (weeks 5-8).** Policy-as-code enforcement on one real use case, in a sidecar, no platform rebuild required.
3. **Rollout.** Extend to production AI workloads, wire to your monitoring and audit pipeline, and keep the posture current through the ongoing incident watch.

The honest one-liner

Australian organisations are deploying agents that can act. The first documented harm has already happened here. SHIELD-AI gives you a way to know, test, and demonstrate what your agents are allowed to do, in terms your regulator will accept.

Security Architecture Practice, Sydney, Australia.

Framework edition: SHIELD-AI v1.8, August 24, 2026