01 · Pillar → Control Atlas

A simplified inventory of the **seven pillars** and their **69 controls**. Use this to see the full control surface at a glance: every control belongs to exactly one pillar, and the pillars are ordered by the logical dependencies they introduce (execution environment is the foundation).

Pillar to control atlas

Pillar to control atlas

Controls by pillar

EE — Execution Environment

ID	Control
EE-06	Fail-closed sandbox operation
EE-02	Infrastructure API access control
EE-01	Execution environment boundary enforcement
EE-08	Evaluation-harness isolation and integrity
EE-03	Resource usage limiting
EE-07	Agent control-plane configuration integrity
EE-04	Evaluation sandbox integrity
EE-05	Host-trusted configuration file injection prevention

MS — Model Security

ID	Control
MS-11	Operational data-stream integrity
MS-10	Agent memory-store integrity
MS-13	Transcript and conversation authenticity

ID	Control
MS-02	Training data integrity
MS-07	Deployment context obfuscation
MS-12	MCP discovery-metadata and shared-cache integrity
MS-09	Perception-fidelity enforcement for human-in-the-loop approvals
MS-05	Model behaviour validation
MS-08	MCP
MS-15	Cross-tenant agent memory and state isolation
MS-01	Model weight protection
MS-06	Model inventory and lineage tracking
MS-03	Supply chain security
MS-14	Self-replication barrier
MS-04	Adversarial robustness

IAM — Identity & Access

ID	Control
IAM-02	Scope-bound delegation
IAM-05	Just-in-time access
IAM-08	Credential scope narrowing
IAM-04	Credential isolation
IAM-06	Cross-session boundary enforcement
IAM-09	Multi-tenant agent platform isolation
IAM-03	Human oversight controls
IAM-07	Transitive delegation chain enforcement
IAM-01	Agent identity management

ID	Control
IAM-10	Taint-tracked least-privilege subagent scoping

DS — Data Security

ID	Control
DS-06	Requirement:** Output validation
DS-07	Requirement:** Agent-originated artifact publication control
DS-02	Inference data protection
DS-01	Data classification and labelling
DS-03	Memory sanitisation
DS-05	Training data governance
DS-04	Data minimisation
DS-08	Requirement:** Destination and emergent-channel egress inventory

TP — Tool & Pipeline Security

ID	Control
TP-08	API gateway infrastructure-isolation mode
TP-10	Agent platform inbound control-plane protection
TP-05	Pipeline integrity
TP-16	Tool-connection and launch-path integrity
TP-01	Tool inventory and registry
TP-03	API gateway security
TP-09	Explicit tool registration — no implicit tool exposure
TP-13	Agent dependency-install provenance verification

ID	Control
TP-11	Agent browser and content-rendering isolation
TP-15	In-process code-execution sandbox sufficiency
TP-14	Skill/plugin payload integrity and instruction-channel isolation
TP-07	Cross-tool dependency analysis
TP-04	Tool output validation
TP-02	Tool access control
TP-06	Dependency scanning
TP-12	Tool-call provenance and dispatch verification

MIR — Monitoring, Logging & IR

ID	Control
MIR-02	Behavioural anomaly detection baseline per-agent behaviour established over training period; alerts on deviation – unusual tool, unusual data access pattern, unusual frequency;
MIR-04	Incident response playbooks AI-specific IR procedures distinct from general IR: prompt injection containment, model rollback, tool access revocation, session invalidation;
MIR-03	Prompt audit log
MIR-05	Mandatory reporting integration
MIR-06	Forensic readiness
MIR-01	Agent activity logging

GRM — Governance & Risk

ID	Control
GRM-02	Requirement:** AI risk register
GRM-03	Ethical use policy
GRM-05	Regulatory mapping and compliance reporting
GRM-01	AI governance body
GRM-04	Requirement:** Third-party AI risk management
GRM-06	Requirement:** Board and executive reporting

Signal for architects

- **EE → MS → IAM → DS → TP → MIR → GRM** is the dependency order: an escaped execution environment invalidates every control layered above it.
- **MS (15) and TP (16)** carry the most controls — model integrity and tool/pipeline security are where the framework concentrates its coverage.