01 · Pillar → Control Atlas

A simplified inventory of the **seven pillars** and their **69 controls**. Use this to see the full control surface at a glance: every control belongs to exactly one pillar, and the pillars are ordered by the logical dependencies they introduce (execution environment is the foundation).

Pillar to control atlas

Pillar to control atlas

Controls by pillar

EE — Execution Environment

ID	Control
EE-04	Evaluation sandbox integrity
EE-01	Execution environment boundary enforcement
EE-02	Infrastructure API access control
EE-05	Host-trusted configuration file injection prevention
EE-08	Evaluation-harness isolation and integrity
EE-07	Agent control-plane configuration integrity
EE-03	Resource usage limiting
EE-06	Fail-closed sandbox operation

MS — Model Security

ID	Control
MS-09	Perception-fidelity enforcement for human-in-the-loop approvals
MS-14	Self-replication barrier

ID	Control
MS-03	Supply chain security
MS-02	Training data integrity
MS-08	MCP
MS-15	Cross-tenant agent memory and state isolation
MS-10	Agent memory-store integrity
MS-05	Model behaviour validation
MS-01	Model weight protection
MS-11	Operational data-stream integrity
MS-12	MCP discovery-metadata and shared-cache integrity
MS-13	Transcript and conversation authenticity
MS-06	Model inventory and lineage tracking
MS-04	Adversarial robustness
MS-07	Deployment context obfuscation

IAM — Identity & Access

ID	Control
IAM-03	Human oversight controls
IAM-02	Scope-bound delegation
IAM-04	Credential isolation
IAM-08	Credential scope narrowing
IAM-05	Just-in-time access
IAM-01	Agent identity management
IAM-06	Cross-session boundary enforcement
IAM-07	Transitive delegation chain enforcement
IAM-09	Multi-tenant agent platform isolation

ID	Control
IAM-10	Taint-tracked least-privilege subagent scoping

DS — Data Security

ID	Control
DS-04	Data minimisation
DS-05	Training data governance
DS-03	Memory sanitisation
DS-07	Requirement:** Agent-originated artifact publication control
DS-01	Data classification and labelling
DS-08	Requirement:** Destination and emergent-channel egress inventory
DS-02	Inference data protection
DS-06	Requirement:** Output validation

TP — Tool & Pipeline Security

ID	Control
TP-03	API gateway security
TP-11	Agent browser and content-rendering isolation
TP-15	In-process code-execution sandbox sufficiency
TP-13	Agent dependency-install provenance verification
TP-10	Agent platform inbound control-plane protection
TP-14	Skill/plugin payload integrity and instruction-channel isolation
TP-05	Pipeline integrity

ID	Control
TP-07	Cross-tool dependency analysis
TP-16	Tool-connection and launch-path integrity
TP-09	Explicit tool registration — no implicit tool exposure
TP-12	Tool-call provenance and dispatch verification
TP-01	Tool inventory and registry
TP-02	Tool access control
TP-06	Dependency scanning
TP-08	API gateway infrastructure-isolation mode
TP-04	Tool output validation

MIR — Monitoring, Logging & IR

ID	Control
MIR-03	Prompt audit log
MIR-02	Behavioural anomaly detection baseline per-agent behaviour established over training period; alerts on deviation – unusual tool, unusual data access pattern, unusual frequency;
MIR-04	Incident response playbooks AI-specific IR procedures distinct from general IR: prompt injection containment, model rollback, tool access revocation, session invalidation;
MIR-05	Mandatory reporting integration
MIR-01	Agent activity logging
MIR-06	Forensic readiness

GRM — Governance & Risk

ID	Control
GRM-05	Regulatory mapping and compliance reporting
GRM-06	Requirement:** Board and executive reporting
GRM-02	Requirement:** AI risk register
GRM-04	Requirement:** Third-party AI risk management
GRM-03	Ethical use policy
GRM-01	AI governance body

Signal for architects

- **EE → MS → IAM → DS → TP → MIR → GRM** is the dependency order: an escaped execution environment invalidates every control layered above it.
- **MS (15) and TP (16)** carry the most controls — model integrity and tool/pipeline security are where the framework concentrates its coverage.