01 · Pillar → Control Atlas

A simplified inventory of the **seven pillars** and their **69 controls**. Use this to see the full control surface at a glance: every control belongs to exactly one pillar, and the pillars are ordered by the logical dependencies they introduce (execution environment is the foundation).

Pillar to control atlas

Pillar to control atlas

Controls by pillar

EE — Execution Environment

ID	Control
EE-01	Execution environment boundary enforcement
EE-02	Infrastructure API access control
EE-03	Resource usage limiting
EE-04	Evaluation sandbox integrity
EE-05	Host-trusted configuration file injection prevention
EE-06	Fail-closed sandbox operation
EE-07	Agent control-plane configuration integrity
EE-08	Evaluation-harness isolation and integrity

MS — Model Security

ID	Control
MS-01	Model weight protection
MS-02	Training data integrity
MS-03	Supply chain security

ID	Control
MS-04	Adversarial robustness
MS-05	Model behaviour validation
MS-06	Model inventory and lineage tracking
MS-07	Deployment context obfuscation
MS-08	MCP
MS-09	Perception-fidelity enforcement for human-in-the-loop approvals
MS-10	Agent memory-store integrity
MS-11	Operational data-stream integrity
MS-12	MCP discovery-metadata and shared-cache integrity
MS-13	Transcript and conversation authenticity
MS-14	Self-replication barrier
MS-15	Cross-tenant agent memory and state isolation

IAM — Identity & Access

ID	Control
IAM-01	Agent identity management
IAM-02	Scope-bound delegation
IAM-03	Human oversight controls
IAM-04	Credential isolation
IAM-05	Just-in-time access
IAM-06	Cross-session boundary enforcement
IAM-07	Transitive delegation chain enforcement
IAM-08	Credential scope narrowing
IAM-09	Multi-tenant agent platform isolation

ID	Control
IAM-10	Taint-tracked least-privilege subagent scoping

DS — Data Security

ID	Control
DS-01	Data classification and labelling
DS-02	Inference data protection
DS-03	Memory sanitisation
DS-04	Data minimisation
DS-05	Training data governance
DS-06	Output validation
DS-07	Agent-originated artifact publication control
DS-08	Destination and emergent-channel egress inventory

TP — Tool & Pipeline Security

ID	Control
TP-01	Tool inventory and registry
TP-02	Tool access control
TP-03	API gateway security
TP-04	Tool output validation
TP-05	Pipeline integrity
TP-06	Dependency scanning
TP-07	Cross-tool dependency analysis
TP-08	API gateway infrastructure-isolation mode
TP-09	Explicit tool registration — no implicit tool exposure
TP-10	Agent platform inbound control-plane protection

ID	Control
TP-11	Agent browser and content-rendering isolation
TP-12	Tool-call provenance and dispatch verification
TP-13	Agent dependency-install provenance verification
TP-14	Skill/plugin payload integrity and instruction-channel isolation
TP-15	In-process code-execution sandbox sufficiency
TP-16	Tool-connection and launch-path integrity

MIR — Monitoring, Logging & IR

ID	Control
MIR-01	Agent activity logging
MIR-02	Behavioural anomaly detection baseline per-agent behaviour established over training period; alerts on deviation – unusual tool, unusual data access pattern, unusual frequency;
MIR-03	Prompt audit log
MIR-04	Incident response playbooks AI-specific IR procedures distinct from general IR: prompt injection containment, model rollback, tool access revocation, session invalidation;
MIR-05	Mandatory reporting integration
MIR-06	Forensic readiness

GRM — Governance & Risk

ID	Control
GRM-01	AI governance body
GRM-02	AI risk register
GRM-03	Ethical use policy
GRM-04	Third-party AI risk management
GRM-05	Regulatory mapping and compliance reporting
GRM-06	Board and executive reporting

Signal for architects

- **EE → MS → IAM → DS → TP → MIR → GRM** is the dependency order: an escaped execution environment invalidates every control layered above it.
- **MS (15) and TP (16)** carry the most controls — model integrity and tool/pipeline security are where the framework concentrates its coverage.