01 · Pillar → Control Atlas

A simplified inventory of the **seven pillars** and their **69 controls**. Use this to see the full control surface at a glance: every control belongs to exactly one pillar, and the pillars are ordered by the logical dependencies they introduce (execution environment is the foundation).

Pillar to control atlas

Pillar to control atlas

Controls by pillar

EE — Execution Environment

ID	Control
EE-03	Resource usage limiting
EE-01	Execution environment boundary enforcement
EE-04	Evaluation sandbox integrity
EE-06	Fail-closed sandbox operation
EE-08	Evaluation-harness isolation and integrity
EE-05	Host-trusted configuration file injection prevention
EE-02	Infrastructure API access control
EE-07	Agent control-plane configuration integrity

MS — Model Security

ID	Control
MS-12	MCP discovery-metadata and shared-cache integrity
MS-02	Training data integrity

ID	Control
MS-07	Deployment context obfuscation
MS-11	Operational data-stream integrity
MS-04	Adversarial robustness
MS-14	Self-replication barrier
MS-03	Supply chain security
MS-05	Model behaviour validation
MS-13	Transcript and conversation authenticity
MS-09	Perception-fidelity enforcement for human-in-the-loop approvals
MS-06	Model inventory and lineage tracking
MS-15	Cross-tenant agent memory and state isolation
MS-08	MCP
MS-10	Agent memory-store integrity
MS-01	Model weight protection

IAM — Identity & Access

ID	Control
IAM-08	Credential scope narrowing
IAM-03	Human oversight controls
IAM-06	Cross-session boundary enforcement
IAM-05	Just-in-time access
IAM-02	Scope-bound delegation
IAM-10	Taint-tracked least-privilege subagent scoping
IAM-01	Agent identity management
IAM-07	Transitive delegation chain enforcement
IAM-09	Multi-tenant agent platform isolation

ID	Control
IAM-04	Credential isolation

DS — Data Security

ID	Control
DS-07	Requirement:** Agent-originated artifact publication control
DS-01	Data classification and labelling
DS-02	Inference data protection
DS-05	Training data governance
DS-04	Data minimisation
DS-03	Memory sanitisation
DS-06	Requirement:** Output validation
DS-08	Requirement:** Destination and emergent-channel egress inventory

TP — Tool & Pipeline Security

ID	Control
TP-01	Tool inventory and registry
TP-03	API gateway security
TP-07	Cross-tool dependency analysis
TP-11	Agent browser and content-rendering isolation
TP-02	Tool access control
TP-10	Agent platform inbound control-plane protection
TP-16	Tool-connection and launch-path integrity
TP-09	Explicit tool registration — no implicit tool exposure

ID	Control
TP-14	Skill/plugin payload integrity and instruction-channel isolation
TP-08	API gateway infrastructure-isolation mode
TP-13	Agent dependency-install provenance verification
TP-05	Pipeline integrity
TP-12	Tool-call provenance and dispatch verification
TP-06	Dependency scanning
TP-04	Tool output validation
TP-15	In-process code-execution sandbox sufficiency

MIR — Monitoring, Logging & IR

ID	Control
MIR-04	Incident response playbooks AI-specific IR procedures distinct from general IR: prompt injection containment, model rollback, tool access revocation, session invalidation;
MIR-06	Forensic readiness
MIR-01	Agent activity logging
MIR-02	Behavioural anomaly detection baseline per-agent behaviour established over training period; alerts on deviation – unusual tool, unusual data access pattern, unusual frequency;
MIR-05	Mandatory reporting integration
MIR-03	Prompt audit log

GRM — Governance & Risk

ID	Control
GRM-01	AI governance body
GRM-04	Requirement:** Third-party AI risk management
GRM-06	Requirement:** Board and executive reporting
GRM-03	Ethical use policy
GRM-02	Requirement:** AI risk register
GRM-05	Regulatory mapping and compliance reporting

Signal for architects

- **EE → MS → IAM → DS → TP → MIR → GRM** is the dependency order: an escaped execution environment invalidates every control layered above it.
- **MS (15) and TP (16)** carry the most controls — model integrity and tool/pipeline security are where the framework concentrates its coverage.