01 · Pillar → Control Atlas

A simplified inventory of the **seven pillars** and their **69 controls**. Use this to see the full control surface at a glance: every control belongs to exactly one pillar, and the pillars are ordered by the logical dependencies they introduce (execution environment is the foundation).

Pillar to control atlas

Pillar to control atlas

Controls by pillar

EE — Execution Environment

ID	Control
EE-08	Evaluation-harness isolation and integrity
EE-03	Resource usage limiting
EE-07	Agent control-plane configuration integrity
EE-02	Infrastructure API access control
EE-01	Execution environment boundary enforcement
EE-05	Host-trusted configuration file injection prevention
EE-04	Evaluation sandbox integrity
EE-06	Fail-closed sandbox operation

MS — Model Security

ID	Control
MS-14	Self-replication barrier
MS-15	Cross-tenant agent memory and state isolation

ID	Control
MS-06	Model inventory and lineage tracking
MS-13	Transcript and conversation authenticity
MS-07	Deployment context obfuscation
MS-04	Adversarial robustness
MS-05	Model behaviour validation
MS-09	Perception-fidelity enforcement for human-in-the-loop approvals
MS-11	Operational data-stream integrity
MS-08	MCP
MS-01	Model weight protection
MS-02	Training data integrity
MS-10	Agent memory-store integrity
MS-03	Supply chain security
MS-12	MCP discovery-metadata and shared-cache integrity

IAM — Identity & Access

ID	Control
IAM-07	Transitive delegation chain enforcement
IAM-10	Taint-tracked least-privilege subagent scoping
IAM-06	Cross-session boundary enforcement
IAM-08	Credential scope narrowing
IAM-03	Human oversight controls
IAM-09	Multi-tenant agent platform isolation
IAM-05	Just-in-time access
IAM-04	Credential isolation
IAM-01	Agent identity management

ID	Control
IAM-02	Scope-bound delegation

DS — Data Security

ID	Control
DS-02	Inference data protection
DS-04	Data minimisation
DS-05	Training data governance
DS-03	Memory sanitisation
DS-08	Requirement:** Destination and emergent-channel egress inventory
DS-07	Requirement:** Agent-originated artifact publication control
DS-06	Requirement:** Output validation
DS-01	Data classification and labelling

TP — Tool & Pipeline Security

ID	Control
TP-04	Tool output validation
TP-10	Agent platform inbound control-plane protection
TP-09	Explicit tool registration — no implicit tool exposure
TP-16	Tool-connection and launch-path integrity
TP-01	Tool inventory and registry
TP-07	Cross-tool dependency analysis
TP-03	API gateway security
TP-06	Dependency scanning
TP-13	

ID	Control
	Agent dependency-install provenance verification
TP-15	In-process code-execution sandbox sufficiency
TP-12	Tool-call provenance and dispatch verification
TP-02	Tool access control
TP-11	Agent browser and content-rendering isolation
TP-08	API gateway infrastructure-isolation mode
TP-05	Pipeline integrity
TP-14	Skill/plugin payload integrity and instruction-channel isolation

MIR — Monitoring, Logging & IR

ID	Control
MIR-04	Incident response playbooks AI-specific IR procedures distinct from general IR: prompt injection containment, model rollback, tool access revocation, session invalidation;
MIR-05	Mandatory reporting integration
MIR-06	Forensic readiness
MIR-03	Prompt audit log
MIR-02	Behavioural anomaly detection baseline per-agent behaviour established over training period; alerts on deviation – unusual tool, unusual data access pattern, unusual frequency;
MIR-01	Agent activity logging

GRM — Governance & Risk

ID	Control
GRM-04	Requirement:** Third-party AI risk management
GRM-02	Requirement:** AI risk register
GRM-03	Ethical use policy
GRM-06	Requirement:** Board and executive reporting
GRM-01	AI governance body
GRM-05	Regulatory mapping and compliance reporting

Signal for architects

- **EE → MS → IAM → DS → TP → MIR → GRM** is the dependency order: an escaped execution environment invalidates every control layered above it.
- **MS (15) and TP (16)** carry the most controls — model integrity and tool/pipeline security are where the framework concentrates its coverage.